

Topic-Specific KGs for Complex Finance Topics

Iain Mackie

University of Glasgow
i.mackie.1@research.gla.ac.uk

Jeffrey Dalton

University of Glasgow
jeff.dalton@glasgow.ac.uk

ABSTRACT

Across the financial domain, researchers answer complex questions by extensively “searching” for relevant information to generate long-form reports or make investment decisions. This paper discusses automating the construction of topic-specific document and entity knowledge graphs (KGs) for complex research topics. We focus on the CODEC dataset, where domain experts (1) create challenging questions, (2) construct long natural language answers, and (3) iteratively search and assess the relevance of documents and entities. For the construction of topic-specific KGs, we show that state-of-the-art ranking systems have headroom for improvement, with specific failings due to lack of context and explicit knowledge representation. Moreover, we demonstrate that entity and document relevance are positively correlated, and that entity-based query feedback improves document ranking effectiveness. We go on to construct topic-specific KGs using retrieval and evaluate using CODEC’s “golden graphs”. Lastly, we point to future work, including adaptive KG retrieval algorithms and GNN-based weighting methods, while highlighting key challenges such as high-quality data, information extraction recall, and the size and sparsity of complex topic graphs.

ACM Reference Format:

Iain Mackie and Jeffrey Dalton. 2022. Topic-Specific KGs for Complex Finance Topics. In *AKBC 2022 Workshop, Knowledge Graphs in Finance and Economics (AKBC Workshop '22)*, November 5, 2021, London, UK. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 INTRODUCTION

Financial analysts research complex topics to identify relevant documents, key entities, and valuable facts. For example, if a researcher wants to understand, *How is the push towards electric cards impacting demand for raw materials?*. In practice, a researcher would use a search engine and reformulate queries to investigate dimensions of the topic, e.g. “Electric vehicles lithium”, “China and electric cards raw materials”, etc. Through this iterative process, the researcher identifies useful sources (i.e. documents) and understands the critical concepts (i.e. entities). Figure 1 shows an example topic-specific KG that analysts would use to support their end goal, such as writing a report or making an investment decision.

A standard approach to automatically construct topic-specific free-text KGs is using retrieval systems to identify relevant documents and entities, and running entity linkers over the documents to extract structured connections [2, 4, 24]. Additionally, neural, sparse, and L2R retrieval methods can provide topic-specific relevance weightings to the KG nodes and edges [4, 24]. This paper focuses primarily on the search component to topic-specific KGs, where in recent years, LLM-based re-ranking [7, 9, 13] and dense retrieval systems [6, 19, 21] ranking leaderboards. However, findings have shown failures on even simple entity queries [17], and headroom for improvement on complex queries where the answers are

long and span multiple documents [10]. Focusing on finance topics within CODEC, we show that current neural models fail due to missing content and a lack of an explicit knowledge representation. We also find real-world finance questions are harder for entity and document ranking systems than for history or politics questions.

Prior work has shown that ad-hoc retrieval [3, 20] and entity ranking [4] can be improved by leveraging entity information. We focus on complex finance topics and use CODEC’s aligned document and entity relevance judgments to understand this relationship better. Specifically, that document relevance positively correlates with the proportion of critical entities contained within the document (+0.21 Pearson Correlation Coefficient) for finance topics. We leverage this relationship to improve initial retrieval with entity query expansion, improving Recall@1000 over BM25 by 22% and 12% of term-based RM3 expansion.

We use strong document and entity retrieval systems to automatically construct topic-specific KGs for CODEC. For evaluation, we use CODEC’s entity and document judgments to develop ground-truth “golden KGs”. This evaluation shows the precision and recall trade-offs when using different run depths and that KGs for complex finance topics are much larger compared to smaller QA-focused graphs [22, 24].

Lastly, we discuss future research directions and challenges. Specifically, we outline opportunities for new adaptive KG retrieval algorithms and GNN-based scoring models. At the same time, we highlight vital challenges such as high-quality data, information extraction recall, and dealing with the size and sparsity of complex topic graphs.

2 CODEC

Complex Document and Entity Collection (CODEC) [11] is a resource that focuses on complex research topics. CODEC covers three tasks: document ranking, entity ranking, and knowledge-grounded answer generation. The document corpus is a 730k-focused web document collection of domain-specific content, the entity KB is Wikipedia [14], and the REL entity linker [18] grounds the documents. CODEC contains essay-style questions across social areas, including finance, history and politics.

In this work, we will focus solely on the finance topics subset that covers wide-ranking research areas, including public sector enterprises, the dot-com bubble, the EU’s monetary union, and the NFT investment category. Domain experts (FX Trader, Investment Banker, Accountant, etc.) are employed to provide deep domain insights for high-quality annotations in this complex domain. This subset includes 14 topics, 127 manual query formulations (over 9 per topic on average), long narrative answers that contain 114 words and 32 entities [15] on average. Furthermore, these topics contain 3,568 entity (254.9 per topic) and 1,970 document relevance judgments (140.7 per topic).

Topic: economics-8

Query: How is the push towards electric cars impacting the demand for raw materials?

Narrative: Mass adoption of electric vehicles (EVs) is expected in the years ahead. This shift has, and will continue to have, a significant impact on demand for specific raw materials. For example, lithium, nickel and cobalt will increase due to battery demand. While oil-based vehicles will decline in the coming decades, which will reduce the need for oil. Any document or entity that discusses the past and future demand shifts of raw materials are relevant.

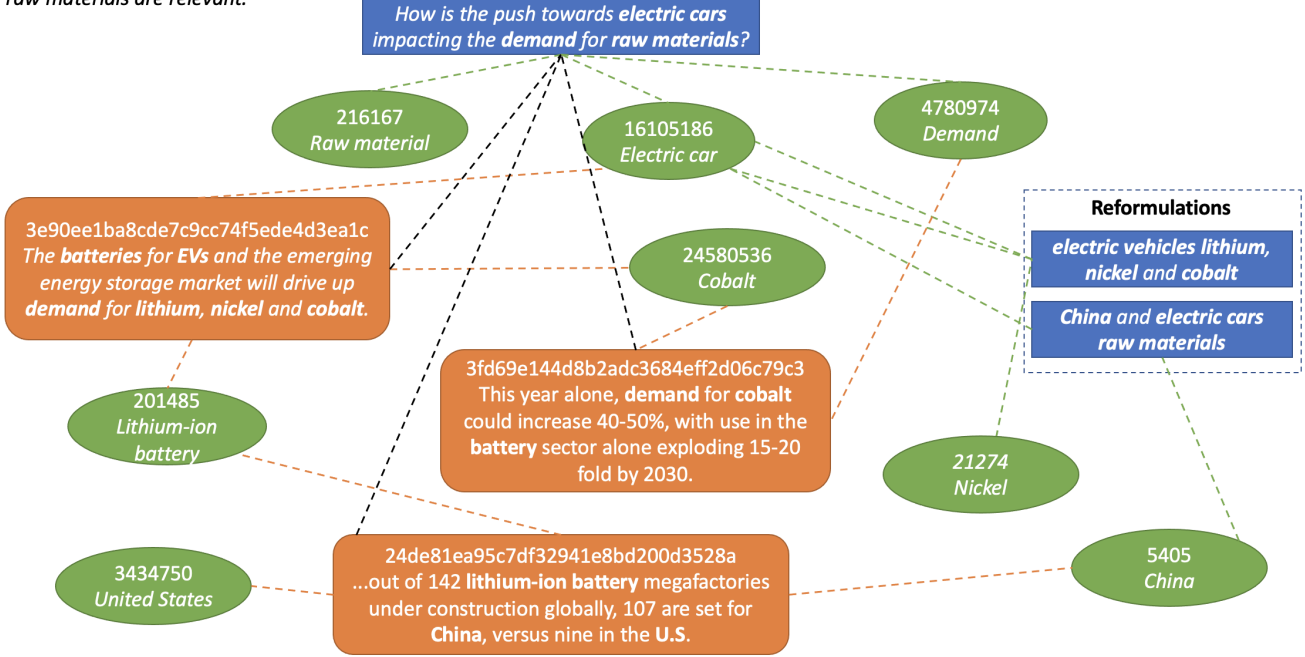


Figure 1: Free-text topic-specific KG constructed for CODEC topic economics-8. Query nodes are in blue, document nodes are in orange, entity nodes are in green, and dotted lines represent entity links.

3 SYSTEMS ANALYSIS

For the task of automatically constructing a topic-specific free-text KG, we focus on the individual components of document and entity ranking. These ranking components are critical as they will directly impact the quality of the KG, which can be explicitly evaluated via relevance judgments. We conduct systems analysis using strong sparse retrieval methods (BM25 [16] and BM25+RM3 [1]), a dense retrieval model (ANCE [21]), and LLM-based re-ranker (T5 [13]).

3.1 Document ranking

Table 1 shows system performance for document ranking. Based on Recall@1000 of 0.735 the best performing method is BM25+RM3, and T5 re-ranking improves MAP to 0.303 and NDCG@100 to 0.487. These measures indicate significant headroom for improvement when compared to other benchmarks [12], where effectiveness will be vital for the accurate construction of topic-specific KGs.

Furthermore, we use BM25+RM3+T5 to compare the hardness of finance topics to other domains present in CODEC (history and politics). Table 2 shows that finance is by far the most challenging domain for both initial retrieval and re-ranking. An example of a hard finance topic is economics-12, *What are the common problems or criticisms aimed at public sector enterprises?*, with Recall@1000 of under 0.55 for all systems. This topic requires a lot of latent knowledge, and analysis of relevant documents shows they contain

Table 1: Document ranking effectiveness on CODEC finance topics.

	NDCG@100	MAP	Recall@1000
BM25	0.330	0.161	0.702
BM25+RM3	0.360	0.182	0.735
ANCE-MaxP	0.349	0.184	0.656
BM25+T5	0.486	0.301	0.702
BM25+RM3+T5	0.487	0.303	0.735
ANCE-MaxP+T5	0.470	0.282	0.656

minimal keyword overlap with the query. Additionally, neural re-ranking systems fail on economic-21 topic, *How much of a threat are ETFs to actively-managed Asset Managers?*, with BM25+RM3+T5 having a MAP of only 0.138. The system runs show that many highly ranked documents do contain entities [ETFs] or [Asset Managers], but current models cannot filter based on the relationships that explain the “threat” ETFs pose to Asset Managers. This motivates more advanced structure reasoning within neural ranking methods.

Lastly, system failures show that low document ranking effectiveness was strongly related to missing documents that contain key topic concepts. For example, all systems had a MAP under 0.150 for topic economics-2, *What technological challenges does Bitcoin*

Table 2: CODEC document ranking effectiveness by domain (system: BM25+RM3+T5).

	NDCG@100	MAP	Recall@1000
History	0.597	0.391	0.850
Politics	0.542	0.344	0.816
Finance	0.487	0.303	0.735

face to becoming a widely used currency?. Highly ranked documents contain anticipated named entities, i.e. [Bitcoin], [Blockchain], and [Satoshi Nakamoto]. However, systems miss the key concepts that are needed to truly understand this information need, i.e. [Transaction time], [Transaction cost], [Quantum technology], [Carbon footprint], and [Cyberattack].

3.2 Entity ranking

Table 3 shows the system performance on the entity ranking task. The performance of entity ranking systems is much lower when compared to document ranking, emphasising that entity ranking is a challenging task within the CODEC setup, i.e. searching for concepts as well as named entities. Neural ranking methods (ANCE and T5) significantly underperform sparse methods (BM25 and BM25+RM3), at least in a zero-shot setting, and we did not include results. The best system is BM25+RM3 with Recall@1000 of 0.651, which is 13% improvement over BM25, showing the clear benefits of pseudo-relevance feedback within entity ranking of complex finance topics.

Table 3: Entity ranking effectiveness on CODEC finance topics.

	NDCG@100	MAP	Recall@1000
BM25	0.325	0.146	0.575
BM25+RM3	0.352	0.177	0.651

Similar to the findings for document ranking, Table 4 supports that finance topics are more challenging for entity ranking systems than history and politics topics. This highlights the complexities of the finance domain, which require contextualisation and reasoning to answer real-world finance questions. Analysing specific topic failures, for example, economics-17 topic “Why is scaling a hardware business more capital intensive than a software business?”. We see that retrieval models cannot identify relevant entities that are categories or sub-categories of [Hardware] or [Software] businesses related to [Capital expenditure], which requires more extensive multi-hop reasoning capabilities.

Table 4: CODEC entity ranking effectiveness by domain (system: BM25+RM3).

	NDCG@100	MAP	Recall@1000
History	0.409	0.224	0.656
Politics	0.415	0.226	0.747
Finance	0.352	0.177	0.651

4 DOCUMENT AND ENTITY RELATIONSHIP

For entity linking, we want high-recall wikification and run WAT [15] entity linker over the entire CODEC 730k document corpus to provide structured connections between documents and entities, enabling construction of a topic-specific graph. This results in 116.5m entity links (159.6 entity links per document), which equated to roughly one entity link every 4 terms. To understand the relationship between document relevance and entity relevance for the finance domain, we connect the document judgments and the relevance of entities mentioned in each document using the entity links and entity judgments. We calculate the Pearson Correlation Coefficient between document relevance and the percentage of entities in the document grouped by entity relevance. We find that entity links (+0.16) support that documents with higher proportions of highly relevant entities are positively correlated with document relevance.

Table 5: Entity-QE document ranking effectiveness against initial retrieval systems.

	NDCG@100	MAP	Recall@1000
BM25	0.330	0.161	0.702
BM25+RM3	0.360	0.182	0.735
Entity-QE	0.404	0.220	0.799

We build on these findings by using an oracle entity feedback method that enriches the query with names of relevant entities taken from entity judgments, Entity-QE. Table 5 shows that Entity-QE improves Recall@1000 over BM25 by 22% and 12% of term-based RM3 expansion. This is an example of how co-optimising both document and entity ranking offers many opportunities to improve the construction of topic-specific KGs.

5 TOPIC-SPECIFIC KGs

We define a simple graph construction approach based on document ranking and entity linking, similar to prior successful query-specific KG approaches [4, 24]. Specifically, given the user’s query Q , we retrieve a set of k_{doc} number of documents D_1, D_2, \dots, D_k . These top-ranked documents are then fed into the WAT entity linking pipeline, which provides entity links $E_1, E_2, \dots, E_n \in D$. Lastly, to reduce the number of non-topic-relevant entities in our KG, we use the entity ranking to filter entities below k_{ent} rank.

Therefore, we create a topic-specific graph $G = (V, E)$, where V are a set of vertices (nodes) of entity, document, or query types. Furthermore, E is a set of edges based on entity links (i.e. document entity mentions) and retrieval edges between the query and documents. The graph is inherently topic-specific as we only contain the more relevant document and entities based on system runs.

5.1 Golden KGs

To theoretically ground the evaluation of graph quality for topic-specific KGs, we define the “golden KGs” for each topic. Doing so will allow us to measure the differences between this ground truth graph and automatically constructed graphs. This evaluation is possible using CODEC as we have ground-truth topic-specific entity

and document judgments. However, we assume WAT document-level entity links are “correct”, which is not always the case, and follow-up work should address this. Specifically, we define the topic “golden graph” as the query connected to all relevant document nodes, and includes entity nodes for document-entity mentions where the entity is itself relevant.

Table 6: Node and edge statistics of CODEC golden KGs.

	Mean Count
Query Nodes	1.00
Document Nodes	44.1
Entity Nodes	25.2
Query-Document Edges	44.1
Document-Entity Edges	1,227.1

Table 6 shows the golden graph mean node and edge statistics across the 14 finance CODEC topics. These golden graphs are relatively small and precision-focused on judged relevant documents and linked relevant entities, with 44.1 document nodes and 25.2 entity nodes on average. The golden graphs on average have 1,227.1 document-to-entity edges, which shows that the golden entities are very central to these topics (48.7 golden edges on average per entity).

When comparing automatically constructed KGs to our golden graphs, we focus on evolution two-fold: (1) precision and (2) recall. Specifically, we define precision as the proportion of nodes or edges present in constructed KGs when comparing them to the golden graphs, $P_{node} = |KG_{golden_nodes}|/|KG_{nodes}|$, this is propositional to works across different graph sizes. This metric focuses on graph scarcity, a significant issue in the accuracy of finance tasks using the KG, for example, factually-grounded text generation or QA. Additionally, we define recall as the metric for the proportion of nodes or edges missing from the constructed KGs when comparing them to the golden graphs, $R_{node} = |KG_{golden_nodes}|/|GoldenGraph_{nodes}|$. In finance, recall measures the quality of missing relevant information, which is vital when use cases require “knowing all the facts”, for example, investment due diligence.

5.2 Experiments

We evaluate automatic KGs using CODECs golden graphs and strong system runs. For example, using the most effective document ranking system, BM25+RM3+T5, we can investigate different KG properties as we vary k_{doc} . We can also use the most effective entity ranking system, BM25+RM3, and varying k_{ent} to vary the automatic exclusion of query and document entity mentions based on relevance. Table 7 and 8 show precision and Recall trade-offs of automatic KG construction as we vary k_{doc} and k_{ent} across [10,50,100,250,1000]. We specifically report results for document nodes, entity nodes and document-entity edges as these will vary graph construction quality.

As expected, when k_{ent} and k_{doc} are at lower thresholds, precision is significantly higher compared to the larger thresholds. For example, when k_{ent} and k_{doc} are equal to 10, the precision of document nodes is 0.56 and entity nodes is 0.62. However, the document-entity edges only have a precision of 0.38, the compounding of errors due to included non-relevant documents and entities.

Table 7: Precision of automatic graph with $k_{doc}, k_{ent} = [10, 50, 100, 250, 1000]$

	10	50	100	250	1000
Document Nodes	0.56	0.33	0.22	0.11	0.03
Entity Nodes	0.62	0.38	0.29	0.17	0.07
Document-Entity Edges	0.38	0.28	0.18	0.09	0.02

Nonetheless, compared to recall of k_{ent} and k_{doc} equal 10, around 90% of nodes and 95 % of golden edges are missing, highlighting the need for large graphs to capture the domain complexities of complex finance topics.

Table 8: Recall of automatic graph with $k_{doc}, k_{ent} = [10, 50, 100, 250, 1000]$

	10	50	100	250	1000
Document Nodes	0.13	0.37	0.50	0.63	0.73
Entity Nodes	0.10	0.32	0.47	0.66	0.92
Document-Entity Edges	0.05	0.23	0.33	0.50	0.67

Around k_{ent} and k_{doc} equal to 50, the precision and recall trends cross each other, with nodes scoring 0.35 and edges scoring 0.25. This highlights the relatively low-performing graph construction, driven by two-fold: (1) this is a challenging task incorporating multiple systems and complex topics, and (2) more advanced methods are required to improve performance. Specifically, future work will focus on incorporating more graph-based pruning techniques, the performance of entity linking techniques [5], and graphs-information being used as part of the retrieval and re-ranking processes.

When k_{ent} and k_{doc} are increased to 1000, recall for documents is 0.73, entities are 0.92, and edges are 0.67. This shows that most of the relevant nodes and edges are within the original retrieval runs and highlights opportunities to improve the neural re-ranking of documents and entities. Nonetheless, it is worth highlighting that entity recall figures are artificially high due to 40% of the relevant entities not being within the golden graphs. Specifically, driven by (1) recall failures of entity linking or (2) some relevant entities are not mentioned in any relevant document.

Overall, this analysis shows that CODEC dataset can be effectively used to benchmark automatic knowledge graph construction on complex finance topics. We show that simple methods can construct KGs and how golden KGs can evaluate construction performance, highlighting specific areas for improvement. Lastly, we show that the automatic graphs constructed for complex finance topics can be very large, with high-recall KGs having 1k+ of nodes and 100k+ edges.

6 FUTURE WORK

We plan to extensively research automated graph construction on complex topics, including a focus on the finance domain due to the domain complexity. Specifically, building upon recent adaptive retrieval work [8], we plan to leverage entities and graph-based information to improve system recall.

Additionally, building on successful applications with factoid and multiple choice QA [22, 24], we aim to develop GNN-based scoring

models for complex research topics, which will leverage both LLM embeddings and graph-based information. However, these examples in QA often only contain a single correct node and tens or low hundreds of additional context nodes. Whereas complex finance topics can contain several queries, a hundred relevant document and entity nodes, and thousands of context nodes. This scale and potential sparsity will be a specific challenge for GNN-based model training.

Other long-term challenges will also require addressing. For example, the performance of information extraction systems has a significant impact on graph quality [5], with specific failing on tail entities [23]. For example, we find that 40% of the relevant entities were not linked to any relevant document. Additionally, high-quality data capturing the “golden facts” is required for more in-depth system evaluation. Currently, we assume a golden edge between all relevant documents and entities mentioned within that document. Therefore, we will run a user study to generate detailed fact-based data.

7 ACKNOWLEDGEMENTS

We would like to acknowledge all the authors on the original CODEC paper (Paul Owoicho, Carlos Gemmell, Sophie Fischer, and Sean MacAvaney) and Ivan Sekulić for on-going collaboration on topic-specific KGs. This work is supported by the 2019 Bloomberg Data Science Research Grant and the Engineering and Physical Sciences Research Council grant EP/V025708/1.

REFERENCES

- [1] Nasreen Abdul-Jaleel, James Allan, W Bruce Croft, Fernando Diaz, Leah Larkey, Xiaoyan Li, Mark D Smucker, and Courtney Wade. 2004. UMass at TREC 2004: Novelty and HARD. *Computer Science Department Faculty Publication Series* (2004), 189.
- [2] Jeffrey Dalton and Laura Dietz. 2013. Constructing query-specific knowledge bases. In *Proceedings of the 2013 workshop on Automated knowledge base construction*. 55–60.
- [3] Jeffrey Dalton, Laura Dietz, and James Allan. 2014. Entity query feature expansion using knowledge base links. *Proceedings of the 37th international ACM SIGIR conference on Research & development in information retrieval* (2014).
- [4] Laura Dietz. 2019. ENT Rank: Retrieving entities for topical information needs through entity-neighbor-text relations. In *Proceedings of the 42nd international ACM SIGIR conference on research and development in information retrieval*. 215–224.
- [5] Amina Kadry and Laura Dietz. 2017. Open relation extraction for support passage retrieval: Merit and open issues. In *Proceedings of the 40th International ACM SIGIR Conference on Research and Development in Information Retrieval*. 1149–1152.
- [6] Omar Khattab and Matei Zaharia. 2020. Colbert: Efficient and effective passage search via contextualized late interaction over bert. In *Proc. of SIGIR*. 39–48.
- [7] Canjia Li, Andrew Yates, Sean MacAvaney, Ben He, and Yingfei Sun. 2020. PARADE: Passage Representation Aggregation for Document Reranking. *arXiv:2008.09093* (2020).
- [8] Sean MacAvaney, Nicola Tonellotto, and Craig Macdonald. 2022. Adaptive Re-Ranking with a Corpus Graph. In *Proceedings of the 31st ACM International Conference on Information & Knowledge Management*. 1491–1500.
- [9] Sean MacAvaney, Andrew Yates, Arman Cohan, and Nazli Goharian. 2019. CEDR: Contextualized Embeddings for Document Ranking. In *Proceedings of the 42nd Annual International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR 2019)*. Paris, France, 1101–1104.
- [10] Iain Mackie, Jeffrey Dalton, and Andrew Yates. 2021. How Deep is your Learning: the DL-HARD Annotated Deep Learning Dataset. In *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval*.
- [11] Iain Mackie, Paul Owoicho, Carlos Gemmell, Sophie Fischer, Sean MacAvaney, and Jeffery Dalton. 2022. CODEC: Complex Document and Entity Collection. In *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '22)*. 3067–3077. <https://doi.org/10.1145/3477495.3531712>
- [12] Tri Nguyen, Mir Rosenberg, Xia Song, Jianfeng Gao, Saurabh Tiwary, Rangan Majumder, and Li Deng. 2016. MS MARCO: A Human Generated Machine Reading Comprehension Dataset. *arXiv:1611.09268v1* (2016).
- [13] Rodrigo Nogueira, Zhiying Jiang, Ronak Pradeep, and Jimmy Lin. 2020. Document Ranking with a Pretrained Sequence-to-Sequence Model. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: Findings*. 708–718.
- [14] Fabio Petroni, Aleksandra Piktus, Angela Fan, Patrick Lewis, Majid Yazdani, Nicola De Cao, James Thorne, Yacine Jernite, Vladimir Karpukhin, Jean Maillard, et al. 2021. KILT: a Benchmark for Knowledge Intensive Language Tasks. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. 2523–2544.
- [15] Francesco Piccinno and Paolo Ferragina. 2014. From TagME to WAT: a new entity annotator. In *Proceedings of the first international workshop on Entity recognition & disambiguation*. 55–62.
- [16] Stephen E Robertson and Steve Walker. 1994. Some simple effective approximations to the 2-poisson model for probabilistic weighted retrieval. In *SIGIR '94*. Springer, 232–241.
- [17] Christopher Sciavolino, Zexuan Zhong, Jinhyuk Lee, and Danqi Chen. 2021. Simple Entity-Centric Questions Challenge Dense Retrievers. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. 6138–6148.
- [18] Johannes M. van Hulst, Faegheh Hasibi, Koen Dercksen, Krisztian Balog, and Arjen P. de Vries. 2020. REL: An Entity Linker Standing on the Shoulders of Giants. In *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '20)*. ACM.
- [19] Ledell Wu, Fabio Petroni, Martin Josifoski, Sebastian Riedel, and Luke Zettlemoyer. 2020. Zero-shot Entity Linking with Dense Entity Retrieval. CoRR abs/1911.03814 (2019). (2020).
- [20] Chenyan Xiong, Jamie Callan, and Tie-Yan Liu. 2017. Word-entity duet representations for document ranking. In *Proceedings of the 40th International ACM SIGIR conference on research and development in information retrieval*. 763–772.
- [21] Lee Xiong, Chenyan Xiong, Ye Li, Kwok-Fung Tang, Jialin Liu, Paul N Bennett, Junaid Ahmed, and Arnold Overwijk. 2020. Approximate Nearest Neighbor Negative Contrastive Learning for Dense Text Retrieval. In *International Conference on Learning Representations*.
- [22] Michihiro Yasunaga, Hongyu Ren, Antoine Bosselut, Percy Liang, and Jure Leskovec. 2021. Qa-gnn: Reasoning with language models and knowledge graphs for question answering. *arXiv preprint arXiv:2104.06378* (2021).
- [23] Ningyu Zhang, Shumin Deng, Zhanlin Sun, Guanying Wang, Xi Chen, Wei Zhang, and Huajun Chen. 2019. Long-tail relation extraction via knowledge graph embeddings and graph convolution networks. *arXiv preprint arXiv:1903.01306* (2019).
- [24] Chen Zhao, Chenyan Xiong, Xin Qian, and Jordan Boyd-Graber. 2020. Complex factoid question answering with a free-text knowledge graph. In *Proceedings of The Web Conference 2020*. 1205–1216.